

21MAT204

Mathematics For Intelligent Systems-3

KALYANA SUNDARAM
CB.EN.U4AIE21120

QUESTION 1

Implement the following SVM versions using CVX for binary datasets (use appropriate binary datasets. Examples: Checkerboard, Spiral, Ring data or any 2 class datasets from net

1. Linear Hard-Margin SVM (L1 SVM)
2. Linear Soft-Margin SVM (L1 SVM)
3. Linear L2 SVM
4. Proximal SVM
5. Non-linear SVM using RBF, Polynomial Kernels
6. Non-linear SVM using RKS

LINEAR HARD-MARGIN SVM (L1 SVM)


- Linear Hard-Margin Support Vector Machines (L1 SVM) is a classification algorithm that tries to find the hyperplane in an N-dimensional space (N — the number of features) that maximally separates the two classes.
- The distance between the hyperplane and the nearest data point from either class is known as the margin.
- Hard-margin SVM can only be used when the data is linearly separable, meaning that it can be separated into two classes by a single straight line.
- One of the main characteristics of hard-margin SVM is that it tries to maximize the margin between the two classes, which makes it more robust to noise.

$$\min_{\mathbf{w}, b} \frac{1}{2} \mathbf{w}^T \mathbf{w}$$
$$\text{s.t.} \quad y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

```

1 - load checkerboard_dataset.mat
2
3 - X = X;
4 - y = y;
5 % Define the optimization problem
6 - cvx_begin
7 -     variables w(2) b xi(size(X,1))
8 -     minimize(norm(xi,1))
9 -     subject to
10 -         y.*(X*w + b) >= 1 - xi
11 -         xi >= 0
12 - cvx_end
13
14 % Extract the weights and bias from the solution
15 - w = w;
16 - b = b;
17 % Classify the data points using the hyperplane equation
18 - predictions = sign(X*w + b);
19
20 % Plot the results
21 - scatter(X(:,1), X(:,2), [], y)
22 - hold on
23 - plot(X(predictions ~= y,1), X(predictions ~= y,2), 'xr')
24 - plot(X(predictions == y,1), X(predictions == y,2), '.g')
25 - xlabel('x1')
26 - ylabel('x2')
27 - legend('Actual Class', 'Misclassified', 'Correctly Classified')
28

```




LINEAR SOFT-MARGIN SVM (L1 SVM)

- Linear Soft-Margin Support Vector Machines (L1 SVM) is a classification algorithm that tries to find the hyperplane in an N-dimensional space (N — the number of features) that maximally separates the two classes while allowing for some misclassification of the data points.
- The distance between the hyperplane and the nearest data point from either class is known as the margin.
- One of the main characteristics of soft-margin SVM is that it can handle non-linearly separable data by allowing for some misclassification.

$$\min \frac{1}{2} ||\mathbf{w}||^2 + C \sum_{i=1}^n \zeta_i$$

$$\text{s.t.} \quad y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \zeta_i \quad \forall i = 1, \dots, n, \quad \zeta_i \geq 0$$

```
l1svm-har-margin.m x l1svm-soft-margin.m x +
1 - load checkerboard_dataset.mat
2
3 - X = X;
4 - y = y;
5 % Set the regularization parameter
6 - lambda = 1;
7
8 % Define the optimization problem
9 - cvx_begin
10 -     variables w(2) b xi(size(X,1))
11 -     minimize(norm(w,2) + lambda*sum(xi))
12 -     subject to
13 -         y.*(X*w + b) >= 1 - xi
14 -         xi >= 0
15 - cvx_end
16
17 % Extract the weights and bias from the solution
18 - w=w;
19 - b=b;
20
21 % Classify the data points using the hyperplane equation
22 - predictions = sign(X*w + b);
23
24 % Plot the results
25 - scatter(X(:,1), X(:,2), [], y)
26 - hold on
27 - plot(X(predictions ~= y,1), X(predictions ~= y,2), 'xr')
28 - plot(X(predictions == y,1), X(predictions == y,2), '.g')
29 - xlabel('x1')
30 - ylabel('x2')
31 - legend('Actual Class', 'Misclassified', 'Correctly Classified')
32
33 %print accuracy
34 - accuracy = sum(predictions == y)/length(y);
35 - fprintf('Accuracy: %f percent \n', accuracy*100);
36
```



Status: Solved
Optimal value (cvx_optval): +4

Accuracy: 50.000000 percent

LINEAR L2 SVM


- L2 support vector machine is a soft-margin SVM which optimizes the sum of squared errors.
- Just as in L1 SVM, the first term in the objective function minimizes the margin, the second term controls the number of misclassifications and C represents the trade-off between the two.
- The only difference between L1 and L2 SVM is that L2 SVM uses the sum of squared slack variables.
- This difference leads to some interesting properties in dual space.

$$\begin{aligned} & \min_{\mathbf{w}, \gamma, \xi} \frac{1}{2} \mathbf{w}^T \mathbf{w} + \frac{C}{2} \sum_{i=1}^m \xi_i^2 \\ & \text{subject to} \quad d_i(\mathbf{w}^T \mathbf{x}_i - \gamma) + \xi_i - 1 \geq 0, \quad 1 \leq i \leq m \\ & \quad \quad \quad \xi_i \geq 0, \quad 1 \leq i \leq m \end{aligned}$$

```

1 - load checkerboard_dataset.mat
2
3 - X = X;
4 - y = y;
5
6 - C = 1;
7
8 % Define the variables
9 - n = size(X,1); % number of samples
10 - d = size(X,2); % number of features
11
12 % Define the objective function and constraints
13 - cvx_begin
14 -     variable w(d)
15 -     variable b
16 -     variable e(n)
17 -     minimize( 0.5 * w'*w + 0.5 * C * sum(e)^2 )
18 -     subject to
19 -         y.*(X*w - b) + e - 1 >= 0
20 - cvx_end
21
22 %plotting the data
23 - figure(1)
24 - plot(X(y==1,1),X(y==1,2),'r+', 'LineWidth',2, 'MarkerSize',7);
25 - hold on
26 - plot(X(y==-1,1),X(y==-1,2), 'bo', 'LineWidth',2, 'MarkerSize',7);
27 - hold on
28 - legend('y=1', 'y=-1')
29 - xlabel('x_1')
30 - ylabel('x_2')
31 - title('Data')
32
33

```



Status: Solved
Optimal value (cvx_optval): +7.66645

PROXIMAL SVM

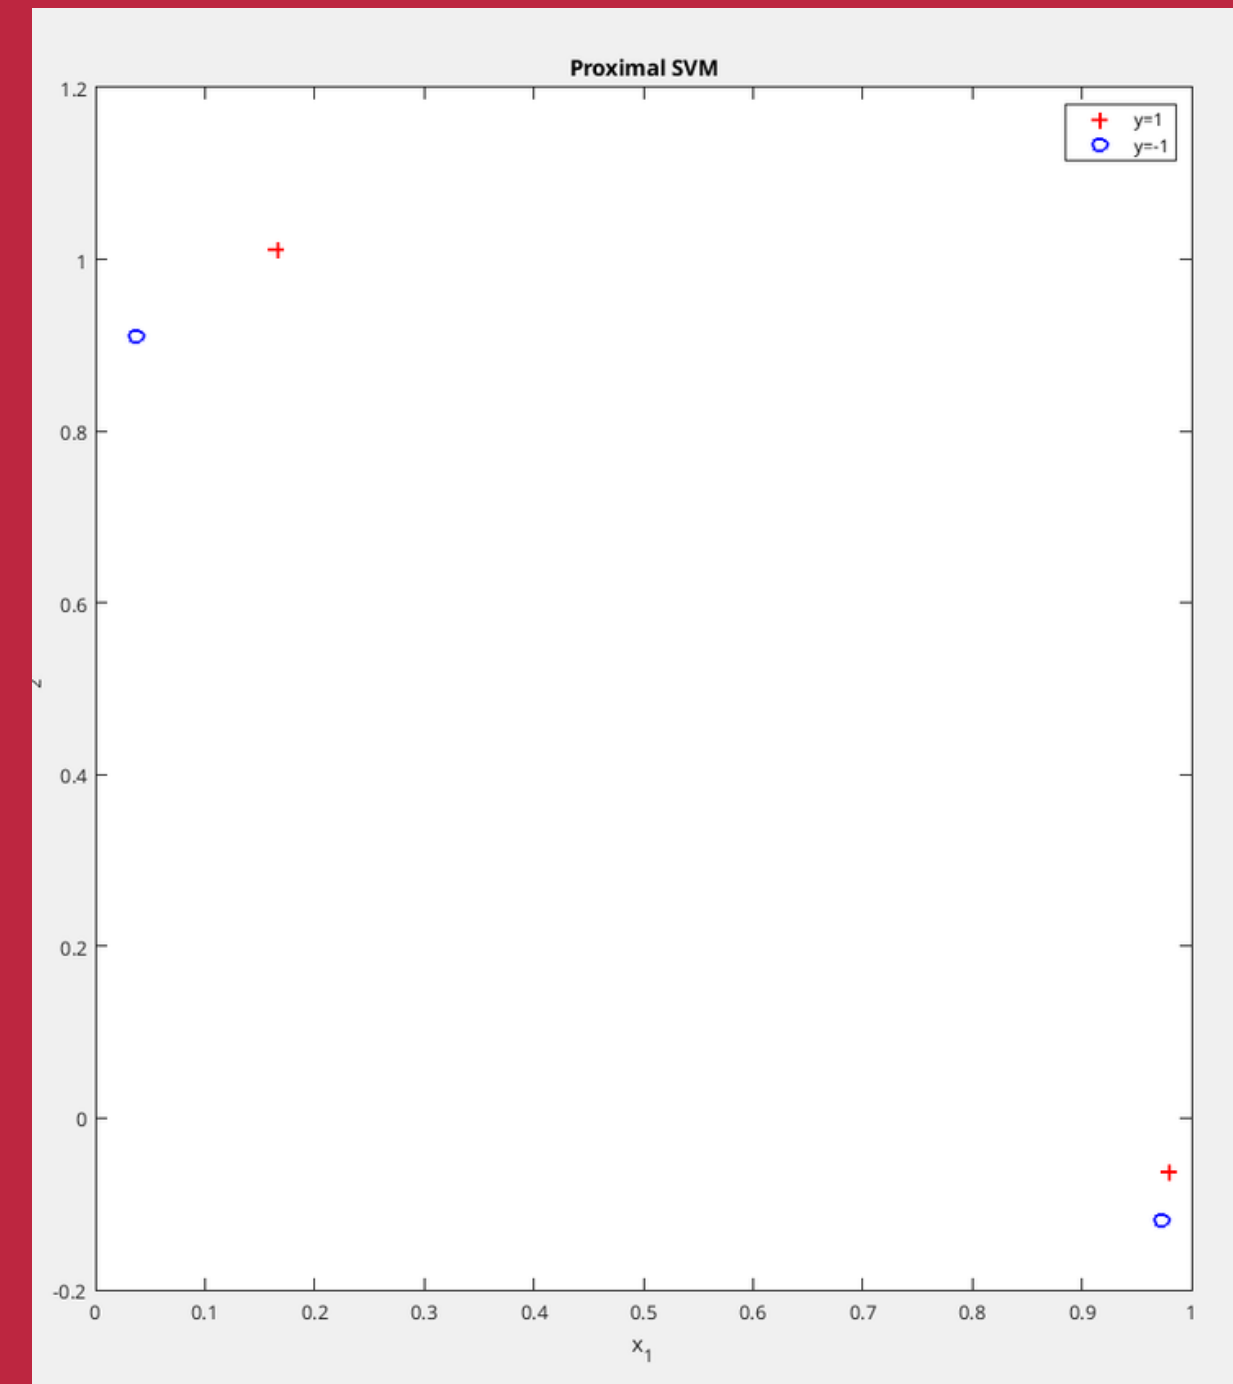
- Instead of a standard support vector machine (SVM) that classifies points by assigning them to one of two disjoint half-spaces, points are classified by assigning them to the closest of two parallel planes (in input or feature space) that are pushed apart as far as possible.

$$\begin{aligned} & \min_{\mathbf{w}, \gamma, \xi} \frac{1}{2} (\mathbf{w}^T \mathbf{w} + \gamma^2) + \frac{C}{2} \sum_{i=1}^m \xi_i^2 \\ & \text{subject to} \quad d_i(\mathbf{w}^T \mathbf{x}_i - \gamma) + \xi_i - 1 = 0, \quad 1 \leq i \leq m \end{aligned}$$

```

l1svm-har-margin.m x l1svm-soft-margin.m x linearl2svm.m x ProximalSVM.m x
1 - load checkerboard_dataset.mat
2
3 - X = X;
4 - y = y;
5 - C = 1;
6
7 % Define the variables
8 - n = size(X,1); % number of samples
9 - d = size(X,2); % number of features
10
11 % Define the objective function and constraints
12 - cvx_begin
13 -     variable w(d)
14 -     variable b
15 -     variable e(n)
16 -     minimize( 0.5 * ( w'*w + b^2 ) + 0.5 * C * sum(e)^2 )
17 -     subject to
18 -         y.*(X*w - b) + e - 1 >= 0
19 - cvx_end
20
21 %plotting the data
22 - figure(1)
23 - plot(X(y==1,1),X(y==1,2),'r+','LineWidth',2,'MarkerSize',7);
24 - hold on
25 - plot(X(y==-1,1),X(y==-1,2),'bo','LineWidth',2,'MarkerSize',7);
26 - hold on
27 - legend('y=1','y=-1')
28 - xlabel('x_1')
29 - ylabel('x_2')
30 - title('Proximal SVM')
31

```



Status: Solved
Optimal value (cvx_optval): +7.66645

NON-LINEAR SVM USING RKS

```
1 - x = rand(2000,1)*5;
2 - y = rand(2000,1)*5;
3 - c = mod((floor(x)+floor(y)),2);
4 - ind = find(c); a = [x(ind),y(ind)];
5 - ind1 = find(c==0);
6 - b = [x(ind1),y(ind1)];
7 - numRows = size(a, 1);
8 - numRows1 = size(b, 1);
9 - A =cat(1, a, b);
10 - di = ones(numRows, 1);
11 - d2 = -1*ones(numRows1, 1);
12 - d = cat(1, di, d2);
13 - numRows2 = size(d, 1);
14 - a = 5;
15 - D = diag(d);
16 - figure;
17 - %gscatter(A(:,1),A(:,2),d,'br','o');
18 - hold on
19 - n = size(A,2);
20 - m = size(A,1);
21 - e = ones(m,1);
22 - c = 10000000;
23 - cvx_begin
24 -     variables w(n) g Psi(m)
25 -     minimize ((0.5*w'*w)+(c*sum(Psi)))
26 -     subject to
27 -         D*(A*w-g*e)+Psi-e >= 0;
28 -         Psi >= 0;
29 - cvx_end
30 - % accuracy
31 - z = sign(A*w-g); r = sum(d==z); Acc = (r/m)*100
32
33 - % plot
34 - x1 = linspace(0,5,100);
35 - y1 = -(w(1)*x1+g)/w(2);
36 - plot(x1,y1,'k-','LineWidth',2)
37 - hold off
38
```

Command Window

Acc =

52.9000

Implement standard LP form using ADMM

1. The first line of code is the declaration of the function `linprog` in matlab. The function `linprog` takes in 6 parameters: the objective function coefficients (`c`), the constraint matrix (`A`), the constraint vector (`b`), the penalty parameter (`rho`), the relaxation parameter (`alpha`) and the return value (`z`). It returns the value of the objective function (`z` and `history`).
2. The second line is the declaration of the variable `t_start`. It is used to record the time when the function is called.
3. The next three lines are the parameters for the algorithm. The first parameter `QUIET` is used to control the output of the algorithm. If `QUIET=0`, the algorithm will output the number of iterations, the objective function value and the norms of the residuals. The second parameter `MAX_ITER` is the maximum number of iterations. The third parameter `ABSTOL` and `RELTOL` are the absolute and relative tolerances for the stopping criteria.
4. The next two lines are the declaration of the variables `m` and `n`. The variable `m` is the number of rows of the constraint matrix. The variable `n` is the number of columns of the constraint matrix.
5. The next three lines are the declaration of the variables `x`, `z` and `u`. The variable `x` is the primal variable. The variable `z` is the dual variable. The variable `u` is the auxiliary variable.
6. The next line is the declaration of the variable `history`. The variable `history` is used to store the objective function value, the norm of the residual, the primal and the dual tolerances for each iteration.
7. The next line is the for loop. The for loop is used to iterate the algorithm. The for loop will execute `MAX_ITER` times or until the stopping criteria is met.

```
+11 q1_1.mlx x q2.mlx x q3.mlx x q4.mlx x DataGen.m x rbf_tst.m x Basis_pursuit.m x
1 function [z, history] = linprog(c, A, b, rho, alpha)
2
3     t_start = tic;
4
5     QUIET    = 0;
6     MAX_ITER = 1000;
7     ABSTOL   = 1e-4;
8     RELTOL   = 1e-2;
9
10    [m, n] = size(A);
11
12
13    x = zeros(n,1);
14    z = zeros(n,1);
15    u = zeros(n,1);
16
17    if ~QUIET
18        fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
19                'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
20    end
21
22    for k = 1:MAX_ITER
23
24        % x-update
25        tmp = [ rho*eye(n), A'; A, zeros(m) ] \ [ rho*(z - u) - c; b ];
26        x = tmp(1:n);
27
28        % z-update with relaxation
29        zold = z;
30        x_hat = alpha*x + (1 - alpha)*zold;
31        z = POS(x_hat + u);
32
33        u = u + (x_hat - z);
34
35        % diagnostics, reporting, termination checks
36
37        history.objval(k) = objective(c, x);
38
39        history.r_norm(k) = norm(x - z);
40        history.s_norm(k) = norm(-rho*(z - zold));
41
42        history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z));
43        history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u);
44
45        if ~QUIET
```

```
randn('state', 0);  
rand('state', 0);  
  
n = 500; % dimension of x  
m = 400; % number of equality constraints  
  
c = rand(n,1) + 0.5; % create nonnegative price vector with mean 1  
x0 = abs(randn(n,1)); % create random solution vector  
  
A = abs(randn(m,n)); % create random, nonnegative matrix A  
b = A*x0;  
  
[x,history] = linprog(c, A, b, 1.0, 1.0);
```

running file

Implement Basis Pursuit using ADMM

```

2 function [z, history] = basis_pursuit(A, b, rho, alpha) %function declaration
3 t_start = tic; %to measure the time taken by the function
4 QUIET = 0; %to print the output or not
5 MAX_ITER = 1000; %maximum number of iterations
6 ABSTOL = 1e-4; %absolute tolerance
7 RELTOL = 1e-2; %relative tolerance
8
9 [m n] = size(A); %size of the matrix A
10 x = zeros(n,1); %initializing x,y,z
11 z = zeros(n,1);
12 u = zeros(n,1);
13
14 if ~QUIET %if QUIET is not 0 then print the following
15     fprintf('%3s\t%10s\t%10s\t%10s\t%10s\t%10s\n', 'iter', ...
16             'r norm', 'eps pri', 's norm', 'eps dual', 'objective');
17 end
18
19 %precompute static variables for x-update (projection on to Ax=b)
20 AAt = A'*A; %AAt is the matrix A multiplied by its transpose
21 P = eye(n) - A' * (AAt \ A); %P is the projection matrix
22 q = A' * (AAt \ b); %q is the projection vector
23
24 for k = 1:MAX_ITER %for loop to iterate for MAX_ITER times or until the termination condition is satisfied
25     % x-update
26     x = P*(z - u) + q;
27
28     % z-update with relaxation
29     zold = z; %zold is the previous value of z
30     x_hat = alpha*x + (1 - alpha)*zold; %x_hat is the relaxation parameter
31     z = shrinkage(x_hat + u, 1/rho); %shrinkage is the soft thresholding function used to update z
32
33     u = u + (x_hat - z); %u is the dual variable which is updated using the relaxation parameter
34
35     %diagnostics, reporting, termination checks
36     history.objval(k) = objective(A, b, x); %objective function is the L1 norm of x
37
38     history.r_norm(k) = norm(x - z); %r_norm is the norm of the difference between x and z
39     history.s_norm(k) = norm(-rho*(z - zold)); %s_norm is the norm of the difference between z and zold
40
41     history.eps_pri(k) = sqrt(n)*ABSTOL + RELTOL*max(norm(x), norm(-z)); %eps_pri is the tolerance for the primal variable
42     history.eps_dual(k) = sqrt(n)*ABSTOL + RELTOL*norm(rho*u); %eps_dual is the tolerance for the dual variable
43
44     if ~QUIET %if QUIET is not 0 then print the following
45         fprintf('%3d\t%10.4f\t%10.4f\t%10.4f\t%10.4f\t%10.2f\n', k, ...
46                 history.r_norm(k), history.eps_pri(k), ...
47                 history.s_norm(k), history.eps_dual(k), history.objval(k));
48     end
49
50     if (history.r_norm(k) < history.eps_pri(k) && ...
51         history.s_norm(k) < history.eps_dual(k))
52         break;
53     end
54 end
55
56 if ~QUIET
57     toc(t_start);
58 end
59 end
60
61 function obj = objective(A, b, x) %function to calculate the objective function
62     obj = norm(x,1);
63 end
64
65 function y = shrinkage(a, kappa) %function to calculate the soft thresholding function
66     y = max(0, a-kappa) - max(0, -a-kappa); %soft thresholding function
67 end
68

```

```

1 clear all;clc;
2 rand('seed', 0);
3 randn('seed', 0);
4
5 n = 30;
6 m = 10;
7 A = randn(m,n);
8
9 x = sprandn(n, 1, 0.1*n);
10 b = A*x;
11
12 xtrue = x;
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29

```

We first import the data from the data file, and then we generate a random vector x with 10% of the entries non-zero. The vector b is generated by multiplying the matrix A and the vector x .

We then call the function `Basis_pursuit.m`, which contains the ADMM algorithm. We set the parameters ρ , α to be 1.0. We then save the output of the function as x , and the history of the algorithm as `history`.

```
[x history] = Basis_pursuit(A, b, 1.0, 1.0);
```

We then plot the objective function value $f(x^k) + g(z^k)$ versus the number of iterations k .

We then plot the residual norm $\|r\|_2$ versus the number of iterations k .

We then plot the residual norm $\|s\|_2$ versus the number of iterations k .

```

K = length(history.objval);

h = figure;
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');



g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');

subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');

```

iter	r norm	eps pri	s norm	eps dual
1	2.1455	0.0240	0.3756	0.0220
2	1.2821	0.0246	1.5779	0.0332
3	0.7697	0.0326	1.2111	0.0354
4	0.8756	0.0371	0.5554	0.0368
5	0.7662	0.0373	0.3617	0.0382
6	0.3543	0.0349	0.5447	0.0387
7	0.2137	0.0334	0.4236	0.0381
8	0.3035	0.0329	0.1683	0.0374
9	0.2479	0.0327	0.0969	0.0371
10	0.1514	0.0328	0.1500	0.0371
11	0.1002	0.0329	0.1479	0.0373
12	0.0870	0.0328	0.1249	0.0375
13	0.0840	0.0327	0.1039	0.0378
14	0.0804	0.0325	0.0895	0.0380
15	0.0774	0.0324	0.0715	0.0381
16	0.0711	0.0324	0.0610	0.0381
17	0.0590	0.0325	0.0522	0.0380
18	0.0513	0.0326	0.0477	0.0381
19	0.0452	0.0327	0.0443	0.0381
20	0.0424	0.0329	0.0394	0.0381
21	0.0426	0.0329	0.0323	0.0381
22	0.0433	0.0329	0.0248	0.0381
23	0.0417	0.0328	0.0214	0.0381
24	0.0373	0.0327	0.0240	0.0382
25	0.0319	0.0326	0.0277	0.0382

Elapsed time is 0.018553 seconds.



Implement LASSO using ADMM

1. The first two lines set the random seed, so that the results are reproducible.
2. We create an m-by-n matrix A, where m is the number of examples and n is the number of features. The matrix is drawn from a standard normal distribution, and the columns are normalized to have unit length.
3. We create a sparse vector x0, where only a fraction p of the entries are non-zero.
4. We create the observation vector b by computing $A \cdot x_0 + \text{noise}$, where noise is drawn from a standard normal distribution and scaled by 0.001.
5. We compute the maximum value of the l1 norm of $A^T b$, which is the largest lambda for which the lasso problem is feasible.
6. We set the regularization parameter lambda to $0.1 \cdot \text{lambda_max}$.
7. We call the Lasso solver. The first argument is the matrix A, the second argument is the vector b, the third argument is the regularization parameter lambda, the fourth argument is the rho parameter, and the fifth argument is the alpha parameter.
8. We plot the objective function value versus the number of iterations.
9. We plot the primal and dual residuals versus the number of iterations.

```

randn('seed', 0);
rand('seed', 0);

m = 1500;      % number of examples
n = 5000;      % number of features
p = 100/n;     % sparsity density

x0 = sprandn(n,1,p);
A = randn(m,n);
A = A*spdiags(1./sqrt(sum(A.^2)),0,n,n); % normalize columns
b = A*x0 + sqrt(0.001)*randn(m,1);

lambda_max = norm(A'*b, 'inf');
lambda = 0.1*lambda_max;

[x history] = Lasso(A, b, lambda, 1.0, 1.0);
K = length(history.objval);



h = figure;
plot(1:K, history.objval, 'k', 'MarkerSize', 10, 'LineWidth', 2);
ylabel('f(x^k) + g(z^k)'); xlabel('iter (k)');

g = figure;
subplot(2,1,1);
semilogy(1:K, max(1e-8, history.r_norm), 'k', ...
    1:K, history.eps_pri, 'k--', 'LineWidth', 2);
ylabel('||r||_2');

subplot(2,1,2);
semilogy(1:K, max(1e-8, history.s_norm), 'k', ...
    1:K, history.eps_dual, 'k--', 'LineWidth', 2);
ylabel('||s||_2'); xlabel('iter (k)');
    
```

iter	r norm	eps pri	s norm	eps dual
1	3.7048	0.0465	0.7250	0.0441
2	2.2654	0.0409	1.7960	0.0653
3	1.0958	0.0529	2.0325	0.0734
4	0.8050	0.0687	1.7219	0.0736
5	0.8619	0.0801	1.2234	0.0704
6	0.8078	0.0864	0.7669	0.0667
7	0.6611	0.0889	0.4398	0.0635
8	0.4906	0.0890	0.2659	0.0612
9	0.3379	0.0878	0.2159	0.0598
10	0.2255	0.0861	0.1987	0.0591
11	0.1585	0.0845	0.1721	0.0590
12	0.1212	0.0833	0.1379	0.0591
13	0.0979	0.0825	0.1044	0.0595
14	0.0799	0.0820	0.0759	0.0598
15	0.0650	0.0819	0.0532	0.0602

Elapsed time is 0.383823 seconds.



THANK YOU